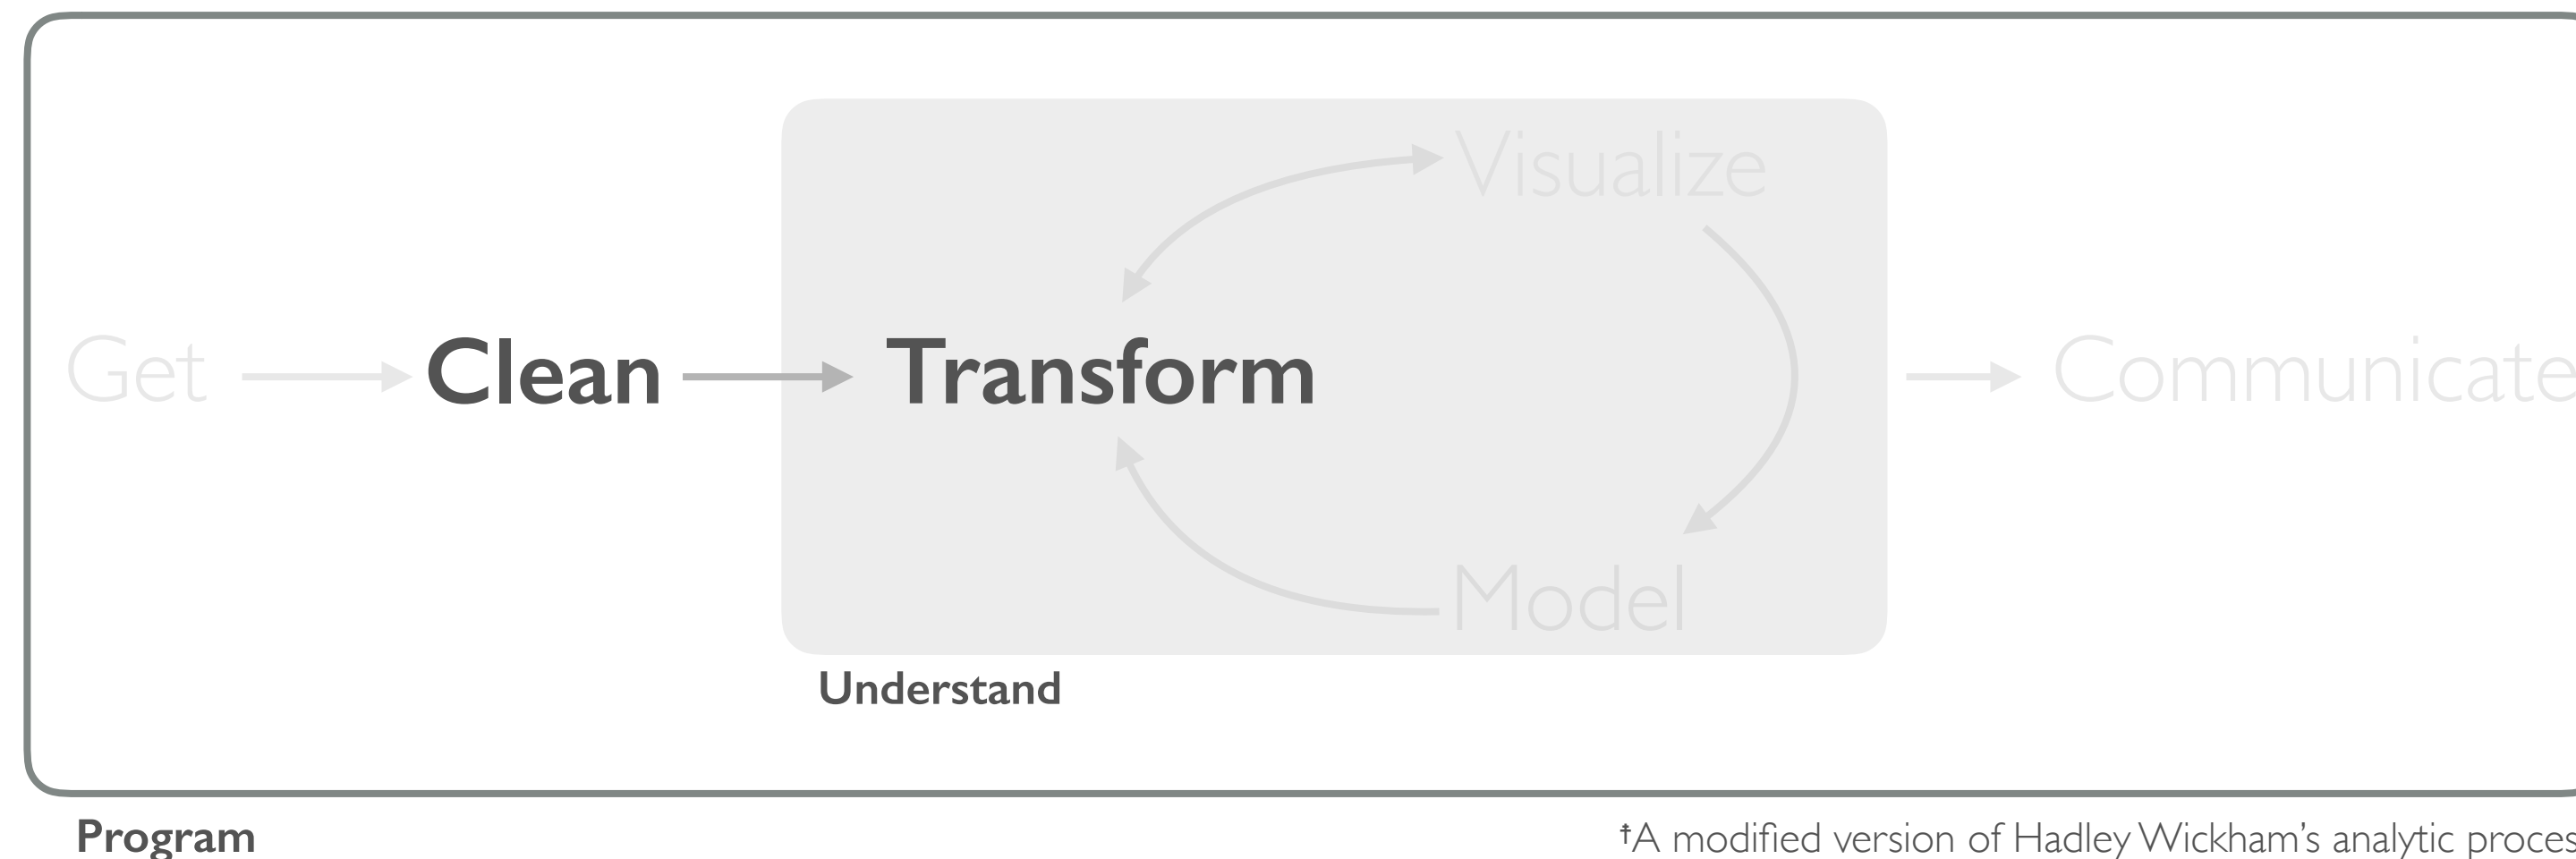


# DATA STRUCTURES



†A modified version of Hadley Wickham's analytic process

# BASICS

## vector

```
0.70 0.86 0.95 0.25 0.52 0.37 0.27 0.80 0.60 0.26
```

## matrix

```
      [,1] [,2] [,3] [,4]  
[1,] 0.70 0.37 0.70 0.37  
[2,] 0.86 0.27 0.86 0.27  
[3,] 0.95 0.80 0.95 0.80  
[4,] 0.25 0.60 0.25 0.60  
[5,] 0.52 0.26 0.52 0.26
```

## data frame

	Sepal.Length	Sepal.Width	Petal.Width	Species
1	5.1	3.5	0.2	setosa
2	4.9	3.0	0.2	setosa
3	4.7	3.2	0.2	setosa
4	4.6	3.1	0.2	setosa
5	5.0	3.6	0.2	setosa
6	5.4	3.9	0.4	setosa
7	4.6	3.4	0.3	setosa
8	5.0	3.4	0.2	setosa
9	4.4	2.9	0.2	setosa
10	4.9	3.1	0.1	setosa

## list

```
$item1  
[1] 1 2 3  
  
$item2  
[1] "a" "b" "c" "d" "e"  
  
$item3  
[1] TRUE FALSE TRUE TRUE  
  
$item4  
      [,1] [,2] [,3]  
[1,] 1    4    7  
[2,] 2    5    8  
[3,] 3    6    9
```

# PREREQUISITES

Re-start your R session

- **Windows:** Ctrl+Shift+F10
- **Mac:** Command+Shift+F10

Reload **nycflights13** library

```
library(nycflights13)
```

# DATA FRAMES

	mpg	cyl	disp	hp	drat	wt	qsec	vs	am	gear	carb
Mazda RX4	21.0	6	160.0	110	3.90	2.620	16.46	0	1	4	4
Mazda RX4 Wag	21.0	6	160.0	110	3.90	2.875	17.02	0	1	4	4
Datsun 710	22.8	4	108.0	93	3.85	2.320	18.61	1	1	4	1
Hornet 4 Drive	21.4	6	258.0	110	3.08	3.215	19.44	1	0	3	1
Hornet Sportabout	18.7	8	360.0	175	3.15	3.440	17.02	0	0	3	2
Valiant	18.1	6	225.0	105	2.76	3.460	20.22	1	0	3	1
Duster 360	14.3	8	360.0	245	3.21	3.570	15.84	0	0	3	4
Merc 240D	24.4	4	146.7	62	3.69	3.190	20.00	1	0	4	2
Merc 230	22.8	4	140.8	95	3.92	3.150	22.90	1	0	4	2
Merc 280	19.2	6	167.6	123	3.92	3.440	18.30	1	0	4	4
Merc 280C	17.8	6	167.6	123	3.92	3.440	18.90	1	0	4	4
Merc 450SE	16.4	8	275.8	180	3.07	4.070	17.40	0	0	3	3
Merc 450SL	17.3	8	275.8	180	3.07	3.730	17.60	0	0	3	3
Merc 450SLC	15.2	8	275.8	180	3.07	3.780	18.00	0	0	3	3
Cadillac Fleetwood	10.4	8	472.0	205	2.93	5.250	17.98	0	0	3	4
Lincoln Continental	10.4	8	460.0	215	3.00	5.424	17.82	0	0	3	4
Chrysler Imperial	14.7	8	440.0	230	3.23	5.345	17.42	0	0	3	4
Fiat 128	32.4	4	78.7	66	4.08	2.200	19.47	1	1	4	1
Honda Civic	30.4	4	75.7	52	4.93	1.615	18.52	1	1	4	2
Toyota Corolla	33.9	4	71.1	65	4.22	1.835	19.90	1	1	4	1
Toyota Corona	21.5	4	120.1	97	3.70	2.465	20.01	1	0	3	1
Dodge Challenger	15.5	8	318.0	150	2.76	3.520	16.87	0	0	3	2
AMC Javelin	15.2	8	304.0	150	3.15	3.435	17.30	0	0	3	2
Camaro Z28	13.3	8	350.0	245	3.73	3.840	15.41	0	0	3	4
Pontiac Firebird	19.2	8	400.0	175	3.08	3.845	17.05	0	0	3	2
Fiat X1-9	27.3	4	79.0	66	4.08	1.935	18.90	1	1	4	1
Porsche 914-2	26.0	4	120.3	91	4.43	2.140	16.70	0	1	5	2
Lotus Europa	30.4	4	95.1	113	3.77	1.513	16.90	1	1	5	2
Ford Pantera L	15.8	8	351.0	264	4.22	3.170	14.50	0	1	5	4

# PROPERTIES

- Spreadsheet style data
- 2 dimensions
  - rows
  - columns
- Can contain heterogenous data
- All columns must be of equal length

	year	month	day	dep_time	carrier	tailnum	dest	time_hour
1	2013	1	1	517	UA	N14228	IAH	2013-01-01 05:00:00
2	2013	1	1	533	UA	N24211	IAH	2013-01-01 05:00:00
3	2013	1	1	542	AA	N619AA	MIA	2013-01-01 05:00:00
4	2013	1	1	544	B6	N804JB	BQN	2013-01-01 05:00:00
5	2013	1	1	554	DL	N668DN	ATL	2013-01-01 06:00:00
6	2013	1	1	554	UA	N39463	ORD	2013-01-01 05:00:00
7	2013	1	1	555	B6	N516JB	FLL	2013-01-01 06:00:00
8	2013	1	1	557	EV	N829AS	IAD	2013-01-01 06:00:00
9	2013	1	1	557	B6	N593JB	MCO	2013-01-01 06:00:00
10	2013	1	1	558	AA	N3ALAA	ORD	2013-01-01 06:00:00
11	2013	1	1	558	B6	N793JB	PBI	2013-01-01 06:00:00
12	2013	1	1	558	B6	N657JB	TPA	2013-01-01 06:00:00
13	2013	1	1	558	UA	N29129	LAX	2013-01-01 06:00:00
14	2013	1	1	558	UA	N53441	SFO	2013-01-01 06:00:00
15	2013	1	1	559	AA	N3DUAA	DFW	2013-01-01 06:00:00
16	2013	1	1	559	B6	N708JB	BOS	2013-01-01 05:00:00
17	2013	1	1	559	UA	N76515	LAS	2013-01-01 06:00:00
18	2013	1	1	600	B6	N595JB	FLL	2013-01-01 06:00:00
19	2013	1	1	600	MQ	N542MQ	ATL	2013-01-01 06:00:00
20	2013	1	1	601	B6	N644JB	PBI	2013-01-01 06:00:00
21	2013	1	1	602	DL	N971DL	MSP	2013-01-01 06:00:00
22	2013	1	1	602	MQ	N730MQ	DTW	2013-01-01 06:00:00
23	2013	1	1	606	AA	N633AA	MIA	2013-01-01 06:00:00

The **flights** data we've been working with is a data frame

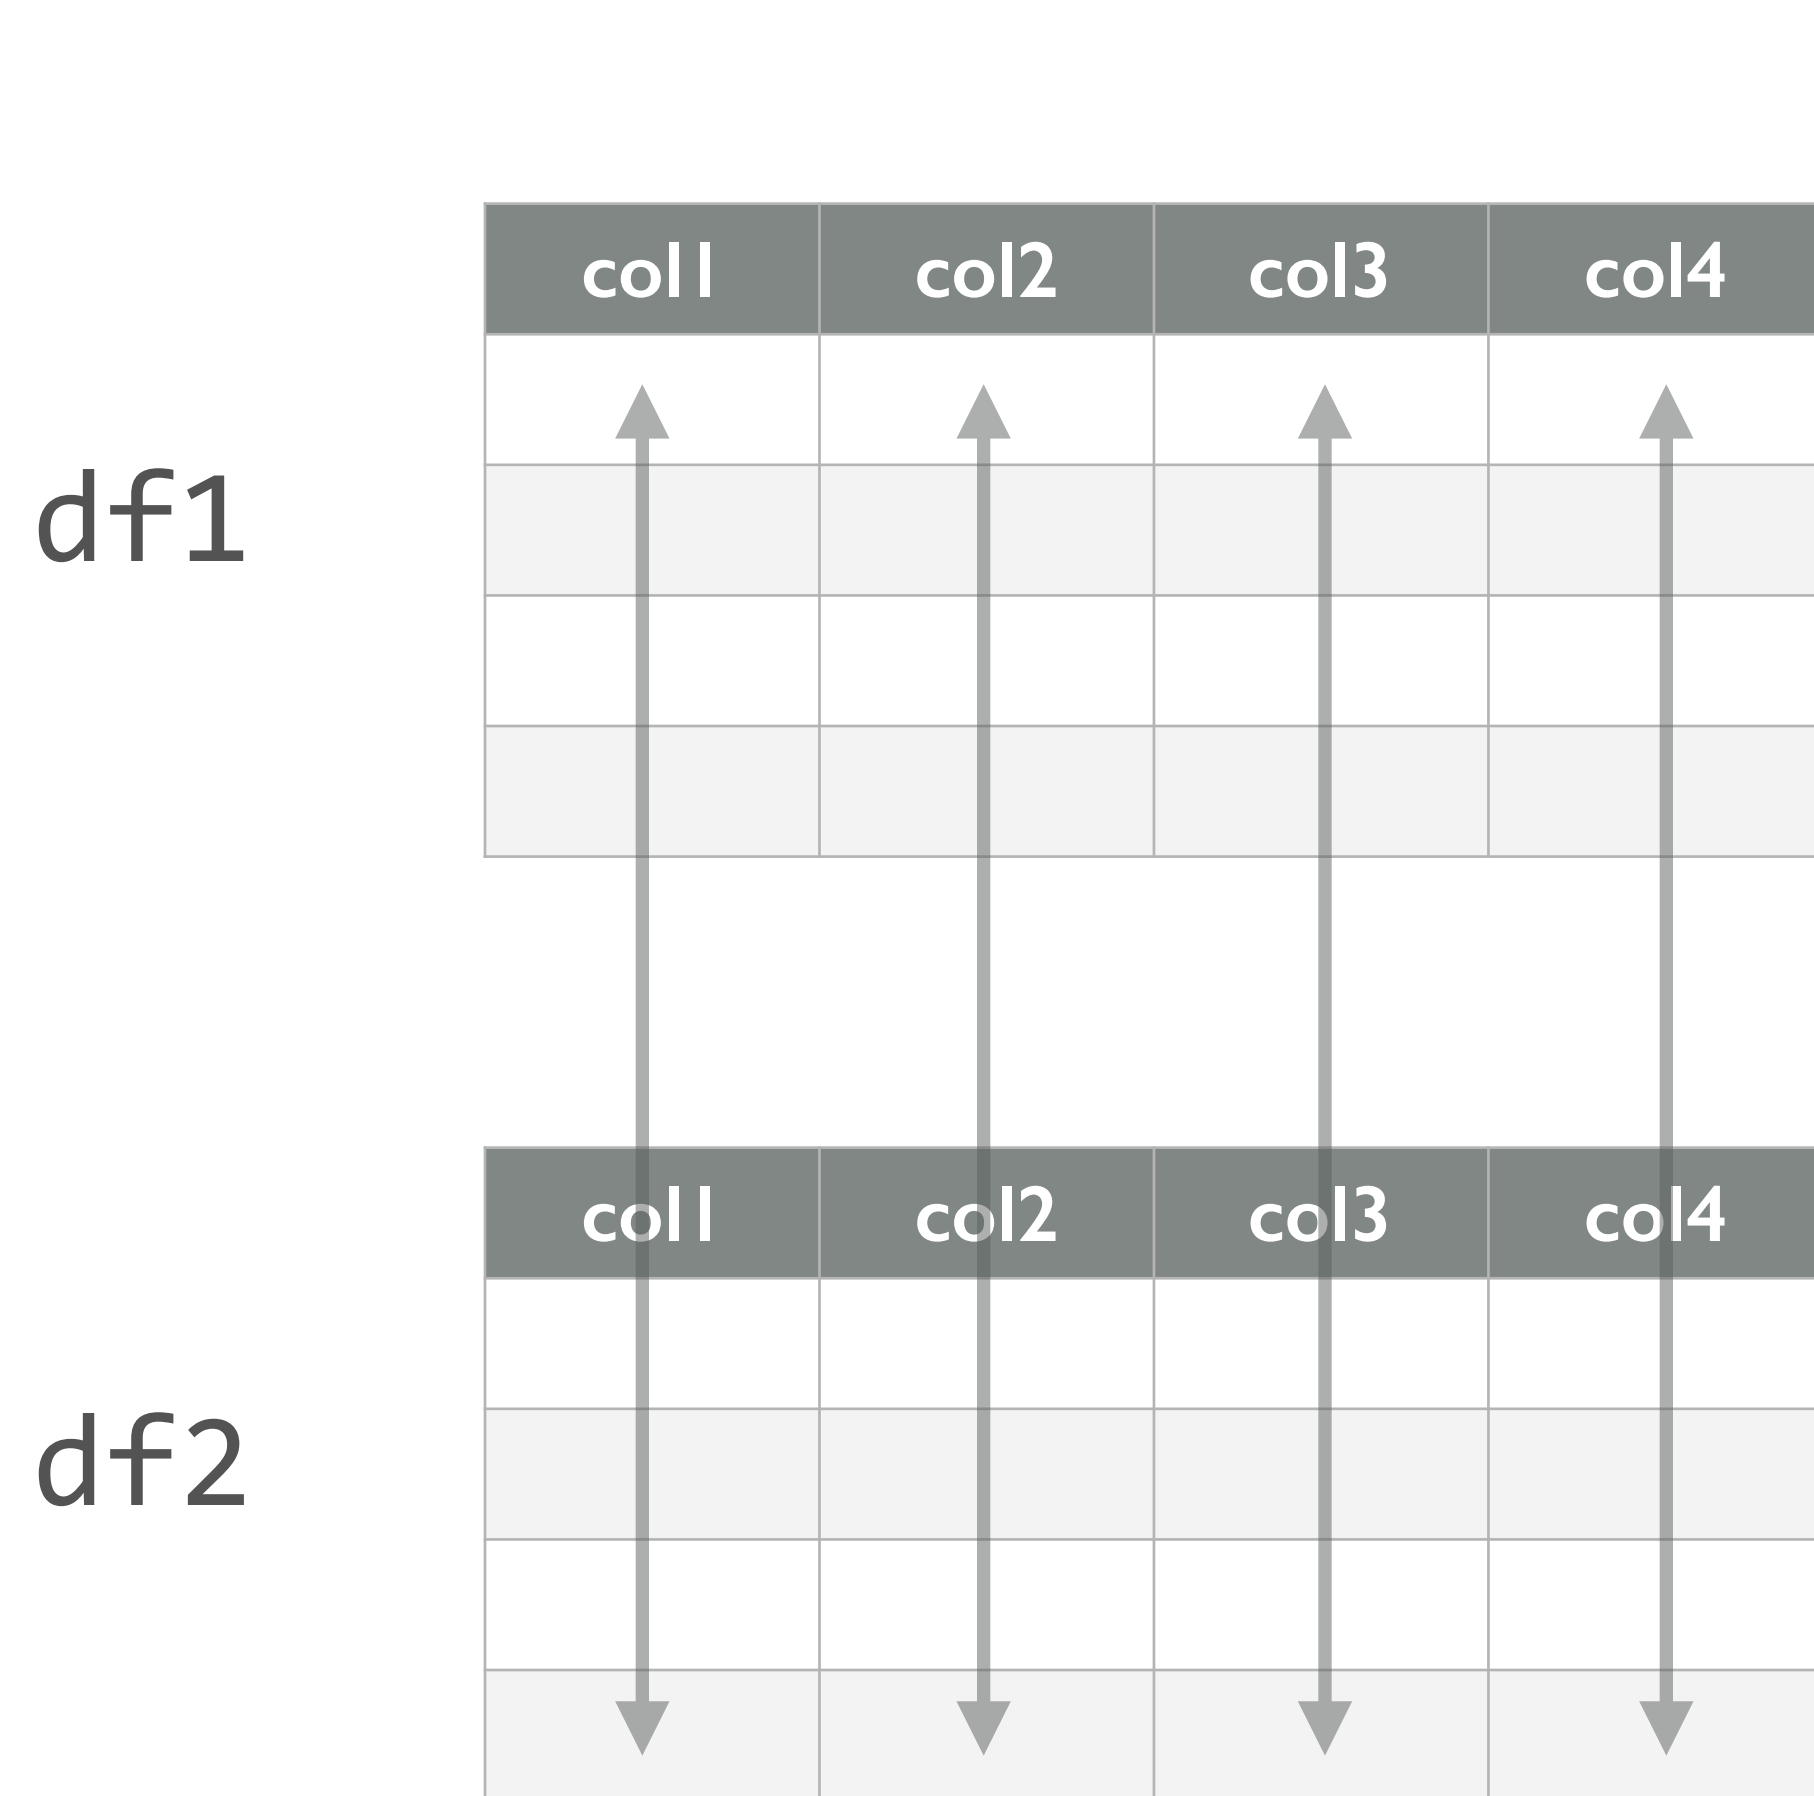
# CREATING

```
df <- data.frame(col1 = 1:3,  
                 col2 = c("this", "is", "text"),  
                 col3 = c(TRUE, FALSE, TRUE),  
                 col4 = c(2.5, 4.2, pi))
```

```
df  
  col1 col2  col3    col4  
1     1 this  TRUE 2.500000  
2     2  is FALSE 4.200000  
3     3 text  TRUE 3.141593
```

*Create this data frame*

# ADDING ON TO

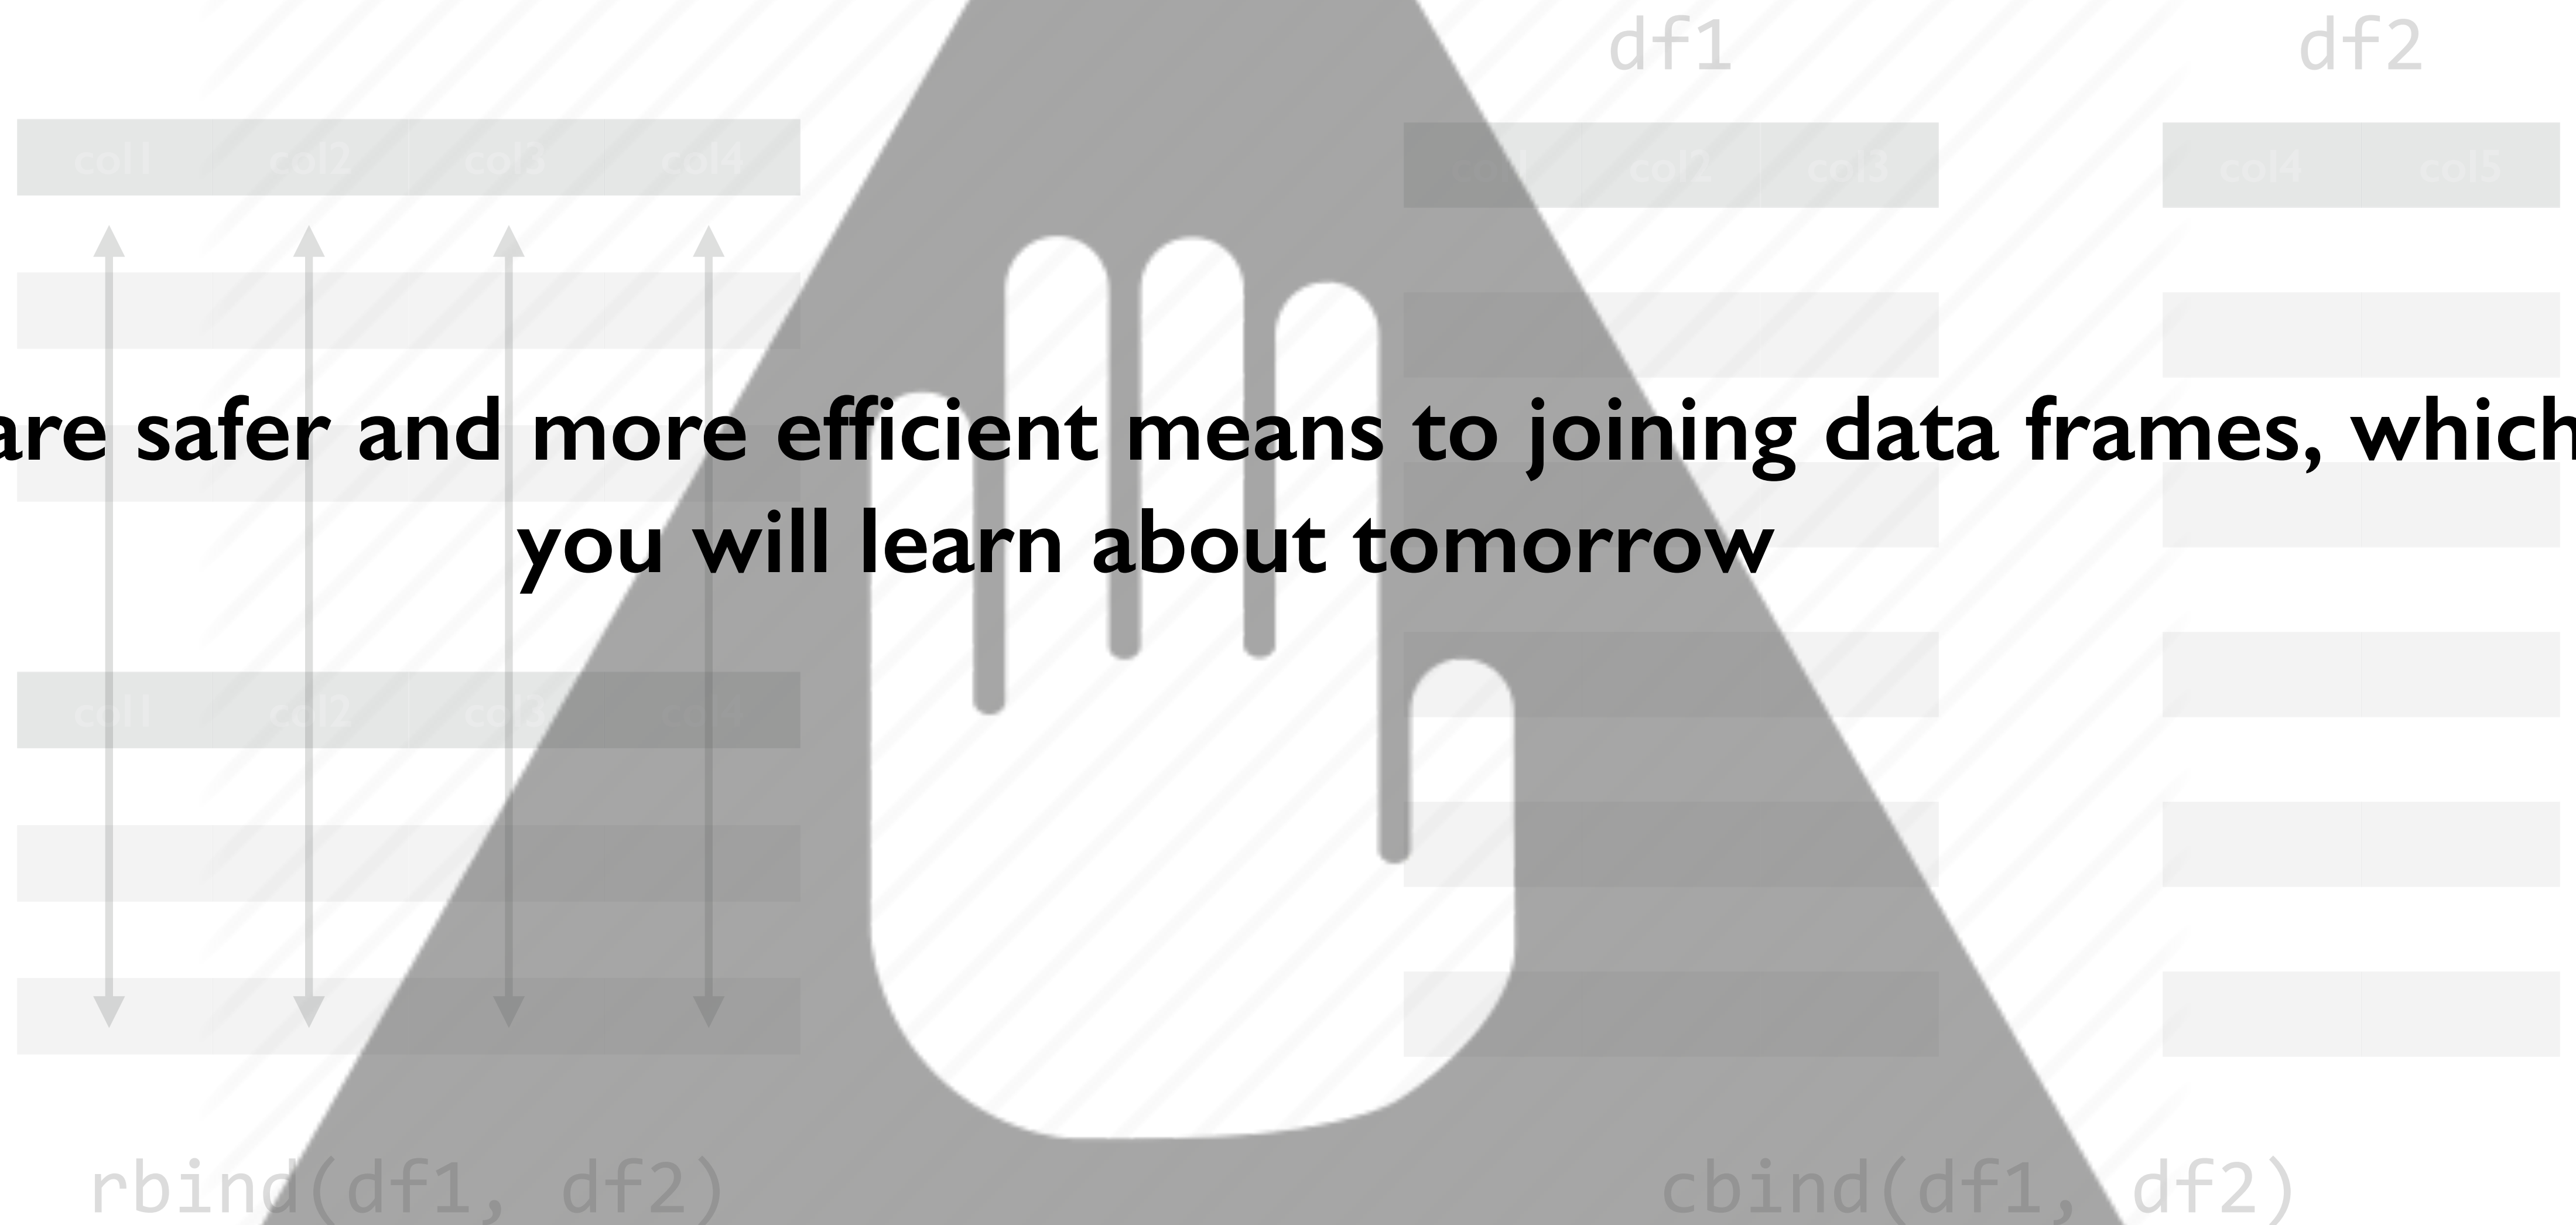


`rbind(df1, df2)`



`cbind(df1, df2)`

# ADDING ON TO





# ATTRIBUTES

Data frames have two main attributes that you care about:

```
# you can also get the number of rows or columns individually with nrow() ncol()
```

```
dim(df)
```

```
[1] 3 4
```

```
names(df)
```

```
[1] "col1" "col2" "col3" "col4"
```

```
names(df) <- c("Col 1", "Col 2", "Col 3", "Col 4")
```

```
df
```

	Col 1	Col 2	Col 3	Col 4
1	1	this	TRUE	2.500000
2	2	is	FALSE	4.200000
3	3	text	TRUE	3.141593

# QUICK SUMMARIES

Get a quick summary of your data frame with `summary()` and `str()`:

For larger data frames you can also use `head(df, n)` and `tail()` to see the first or last  $n$  rows.

```
summary(df)
```

Col 1	Col 2	Col 3	Col 4
Min. :1.0	is :1	Mode :logical	Min. :2.500
1st Qu.:1.5	text:1	FALSE:1	1st Qu.:2.821
Median :2.0	this:1	TRUE :2	Median :3.142
Mean :2.0		NA's :0	Mean :3.281
3rd Qu.:2.5			3rd Qu.:3.671
Max. :3.0			Max. :4.200

```
str(df)
```

```
'data.frame': 3 obs. of 4 variables:
 $ Col 1: int 1 2 3
 $ Col 2: Factor w/ 3 levels "is","text","this": 3 1 2
 $ Col 3: logi TRUE FALSE TRUE
 $ Col 4: num 2.5 2.821 3.142
```

# INDEXING/SUBSETTING

Most of our indexing and subsetting of data frames will be done with **dplyr** functions (**filter** and **select**)

But as you'll see, understanding the `[ ]` functionality is important.

# INDEXING/SUBSETTING

`data.frame[row, col]`

Try these different forms of indexing & subsetting:

```
# extract the second column and all rows using column indexing or the name
df[, 2]
df[, "Col 2"]
```

```
# extract all rows and columns 1 through 3
df[, 1:3]
df[, c("Col 1", "Col 2", "Col 3")]
```

```
# index for first row and all columns
df[1, ]
```

```
# subset for rows
subset(df, `Col 3` == TRUE)
subset(df, `Col 3` == TRUE & `Col 4` > 3, c(2, 4))
```

# YOUR TURN!

1. Using `[ ]`, select the first 1000 rows and the following columns: `month`, `dep_delay`, `carrier`, `distance`, `time_hour`. Save this as `small_flights`.
2. Look at the structure and summary of `small_flights`
3. Rename the columns of `small_flights` to `c("Month", "Delay", "Carrier", "Distance", "Date-Time")`
4. Look at the first and last 15 rows

# SOLUTION

```
# 1
small_flights <- flights[1:1000, c("month", "dep_delay", "carrier", "distance", "time_hour")]

# 2
str(small_flights)
summary(small_flights)

# 3
names(small_flights) <- c("Month", "Delay", "Carrier", "Distance", "Date-Time")

# 4
head(small_flights, 15)
tail(small_flights, 15)
```

# TIBBLES

```
# A tibble: 336,776 × 19
  year month   day dep_time sched_dep_time dep_delay
  <int> <int> <int>   <int>         <int>         <dbl>
1  2013     1     1     517             515           2
2  2013     1     1     533             529           4
3  2013     1     1     542             540           2
4  2013     1     1     544             545          -1
5  2013     1     1     554             600          -6
6  2013     1     1     554             558          -4
7  2013     1     1     555             600          -5
8  2013     1     1     557             600          -3
9  2013     1     1     557             600          -3
10 2013     1     1     558             600          -2
# ... with 336,766 more rows, and 13 more variables:
#   arr_time <int>, sched_arr_time <int>, arr_delay <dbl>,
#   carrier <chr>, flight <int>, tailnum <chr>, origin <chr>,
#   dest <chr>, air_time <dbl>, distance <dbl>, hour <dbl>,
#   minute <dbl>, time_hour <dtm>
```

# PROPERTIES

*Same as data frames but with minor tweaks to make life easier*

*Compare the outputs by running this code:*

```
flights
```

```
tibble::as_tibble(flights)
```

***What differences do you notice?***



# PROPERTIES

A tibble is a data frame with a better printing structure

```
tibble::as_tibble(flights)
# A tibble: 336,776 × 19
   year month   day dep_time sched_dep_time dep_delay arr_time sched_arr_time arr_delay
   <int> <int> <int>   <int>         <int>         <dbl>   <int>         <int>         <dbl>
1  2013     1     1     517           515           2     830           819           11
2  2013     1     1     533           529           4     850           830           20
3  2013     1     1     542           540           2     923           850           33
4  2013     1     1     544           545          -1    1004          1022          -18
5  2013     1     1     554           600          -6     812           837          -25
6  2013     1     1     554           558          -4     740           728           12
7  2013     1     1     555           600          -5     913           854           19
8  2013     1     1     557           600          -3     709           723          -14
9  2013     1     1     557           600          -3     838           846           -8
10 2013     1     1     558           600          -2     753           745            8
# ... with 336,766 more rows, and 10 more variables: carrier <chr>, flight <int>,
#   tailnum <chr>, origin <chr>, dest <chr>, air_time <dbl>, distance <dbl>, hour <dbl>,
#   minute <dbl>, time_hour <dtm>
```

# CREATE

- tibbles are provided by the **tibble** package which is also provided by the **tidyverse** package
- to convert data frames to tibbles just apply **as\_tibble()**

```
library(tidyverse)
as_tibble(iris)
# A tibble: 150 × 5
   Sepal.Length Sepal.Width Petal.Length Petal.Width Species
      <dbl>         <dbl>        <dbl>         <dbl>    <fctr>
1         5.1         3.5          1.4          0.2    setosa
2         4.9         3.0          1.4          0.2    setosa
3         4.7         3.2          1.3          0.2    setosa
4         4.6         3.1          1.5          0.2    setosa
5         5.0         3.6          1.4          0.2    setosa
6         5.4         3.9          1.7          0.4    setosa
7         4.6         3.4          1.4          0.3    setosa
8         5.0         3.4          1.5          0.2    setosa
9         4.4         2.9          1.4          0.2    setosa
10        4.9         3.1          1.5          0.1    setosa
# ... with 140 more rows
```

# PROPERTIES

*All indexing, subsetting, and attribute functions apply to tibbles just as they do to data frames.*

# MATRICES

	[,1]	[,2]	[,3]	[,4]	[,5]	[,6]
[1,]	0.34	0.96	0.36	0.95	0.50	0.98
[2,]	0.47	0.25	0.68	0.65	0.37	0.53
[3,]	0.35	0.93	0.60	0.65	0.14	0.71
[4,]	0.89	0.68	0.07	0.10	0.46	0.20
[5,]	0.28	0.25	0.70	0.36	0.59	0.26
[6,]	0.96	0.42	0.93	0.62	0.24	0.82
[7,]	0.72	0.13	0.47	0.93	0.05	0.23
[8,]	0.82	0.32	0.70	0.84	0.66	0.70
[9,]	0.68	0.04	0.06	0.82	0.78	0.84
[10,]	0.13	0.14	0.46	0.91	0.29	0.82
[11,]	0.45	0.29	0.04	0.12	0.92	0.57
[12,]	0.90	0.81	0.74	0.83	0.91	0.29
[13,]	0.89	0.40	0.71	0.12	0.73	0.08
[14,]	0.05	0.52	0.47	0.53	0.53	0.96
[15,]	0.16	0.59	0.43	0.19	0.37	0.54

# PROPERTIES

- 2 dimensions
  - rows
  - columns
- Can only contain homogenous data
- All columns must be of equal length

	[,1]	[,2]	[,3]	[,4]	[,5]	[,6]
[1,]	0.34	0.96	0.36	0.95	0.50	0.98
[2,]	0.47	0.25	0.68	0.65	0.37	0.53
[3,]	0.35	0.93	0.60	0.65	0.14	0.71
[4,]	0.89	0.68	0.07	0.10	0.46	0.20
[5,]	0.28	0.25	0.70	0.36	0.59	0.26
[6,]	0.96	0.42	0.93	0.62	0.24	0.82
[7,]	0.72	0.13	0.47	0.93	0.05	0.23
[8,]	0.82	0.32	0.70	0.84	0.66	0.70
[9,]	0.68	0.04	0.06	0.82	0.78	0.84
[10,]	0.13	0.14	0.46	0.91	0.29	0.82
[11,]	0.45	0.29	0.04	0.12	0.92	0.57
[12,]	0.90	0.81	0.74	0.83	0.91	0.29
[13,]	0.89	0.40	0.71	0.12	0.73	0.08
[14,]	0.05	0.52	0.47	0.53	0.53	0.96
[15,]	0.16	0.59	0.43	0.19	0.37	0.54

# CREATING

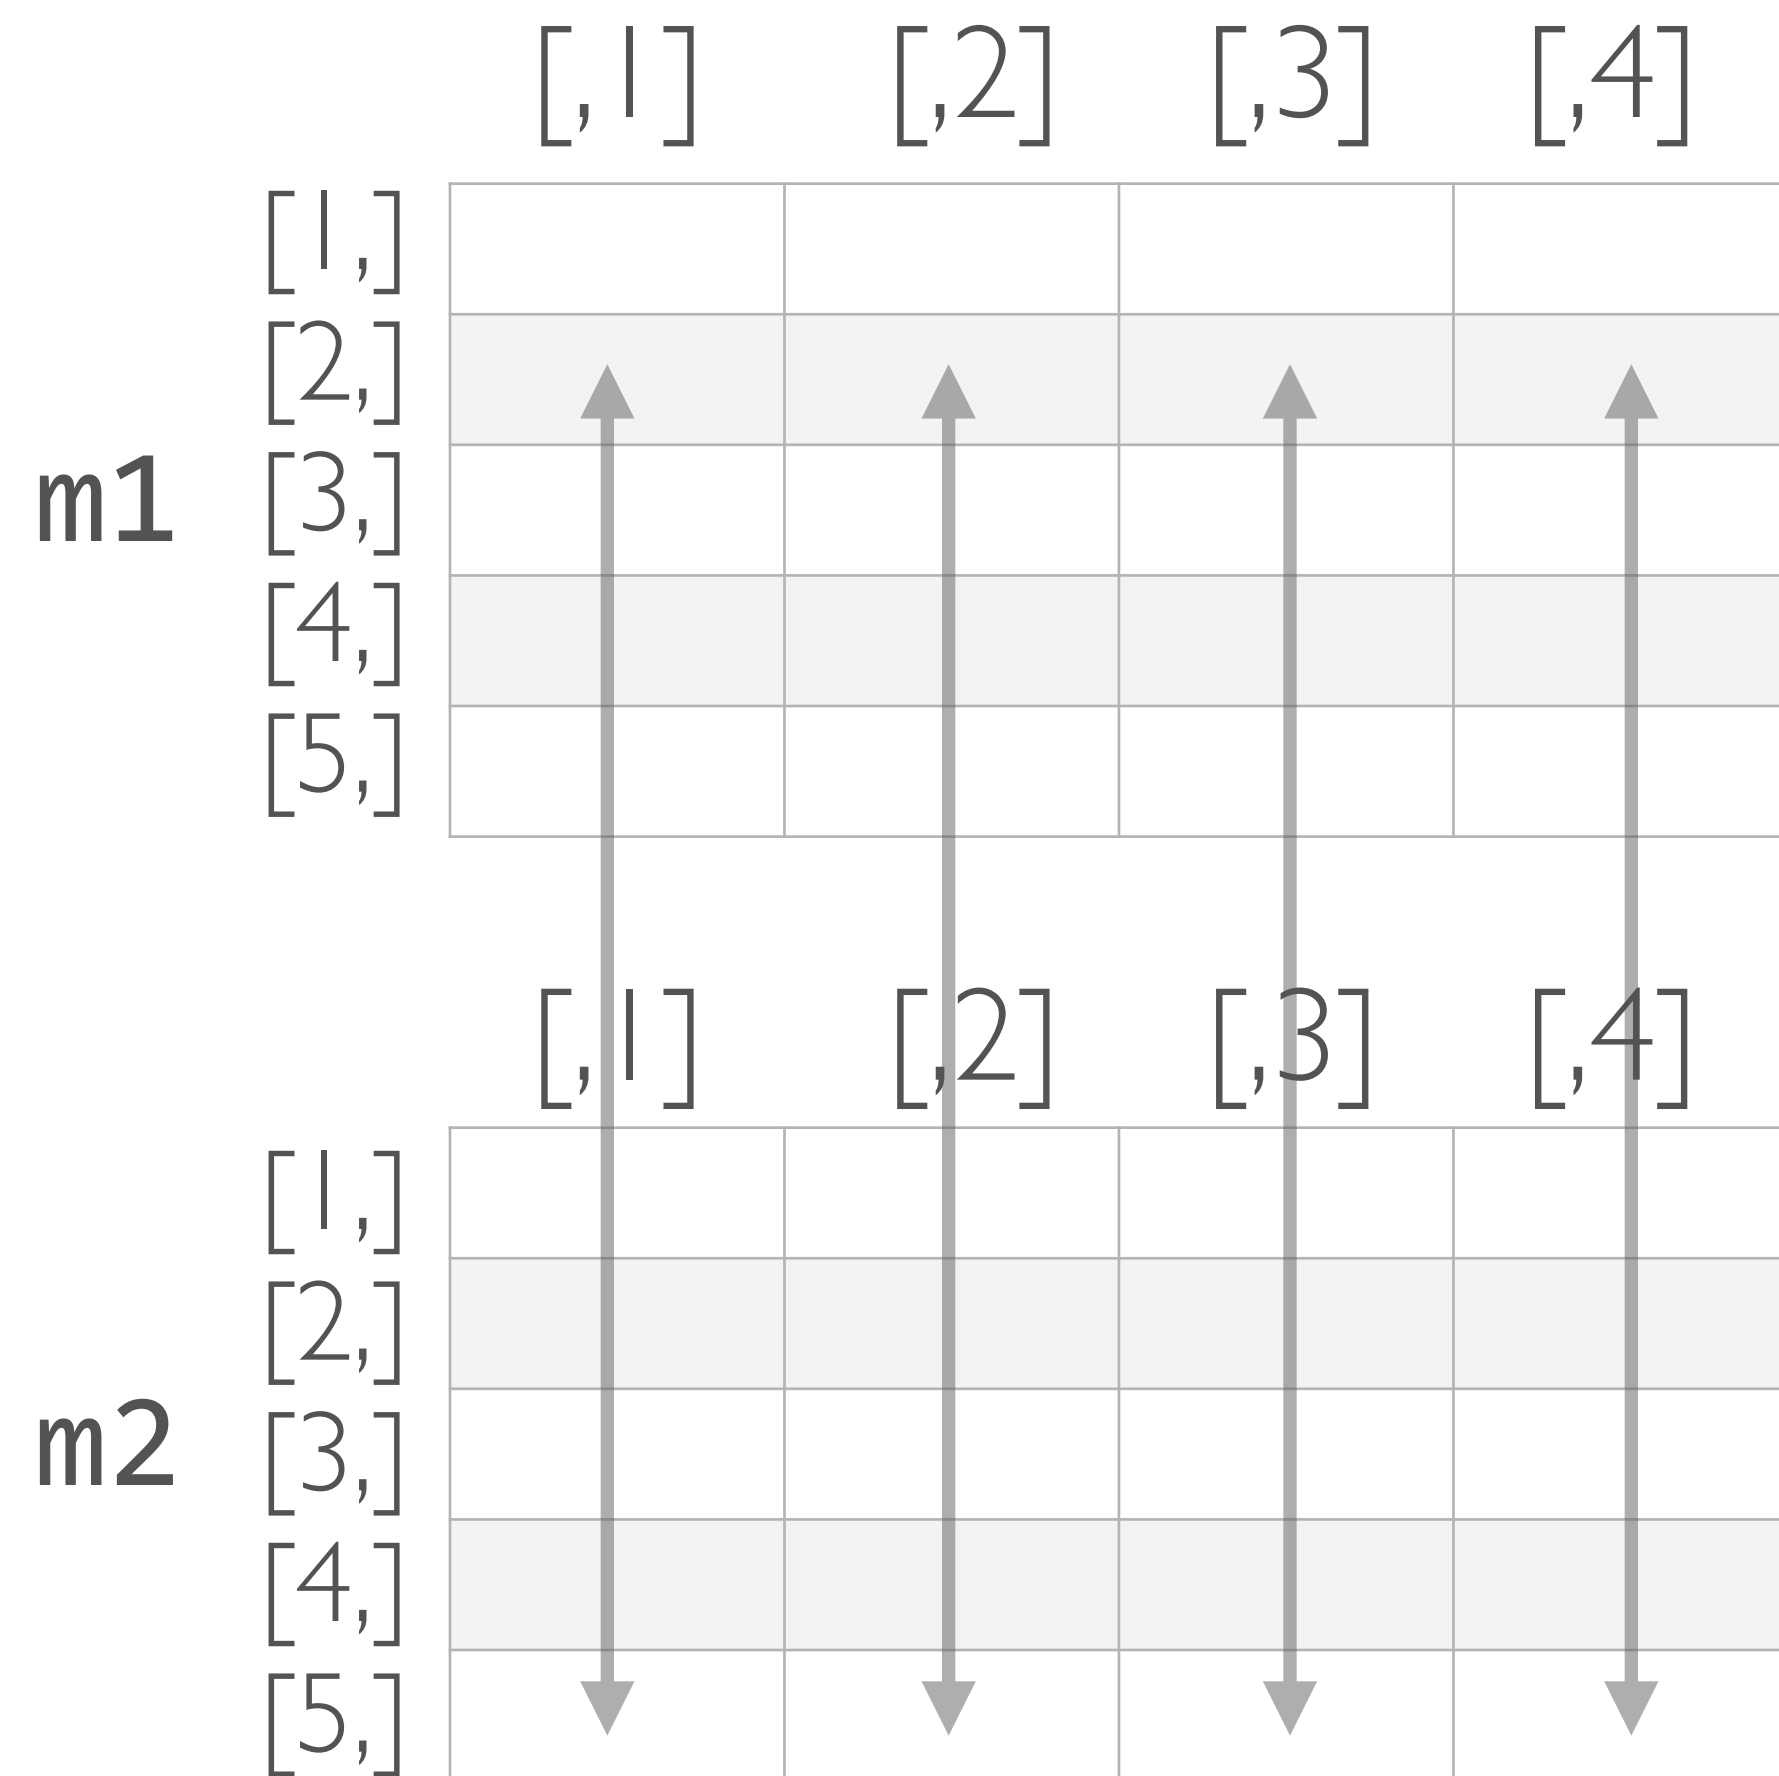
```
set.seed(123)
v1 <- sample(1:10, 25, replace = TRUE)
m1 <- matrix(v1, nrow = 5)
```

m1

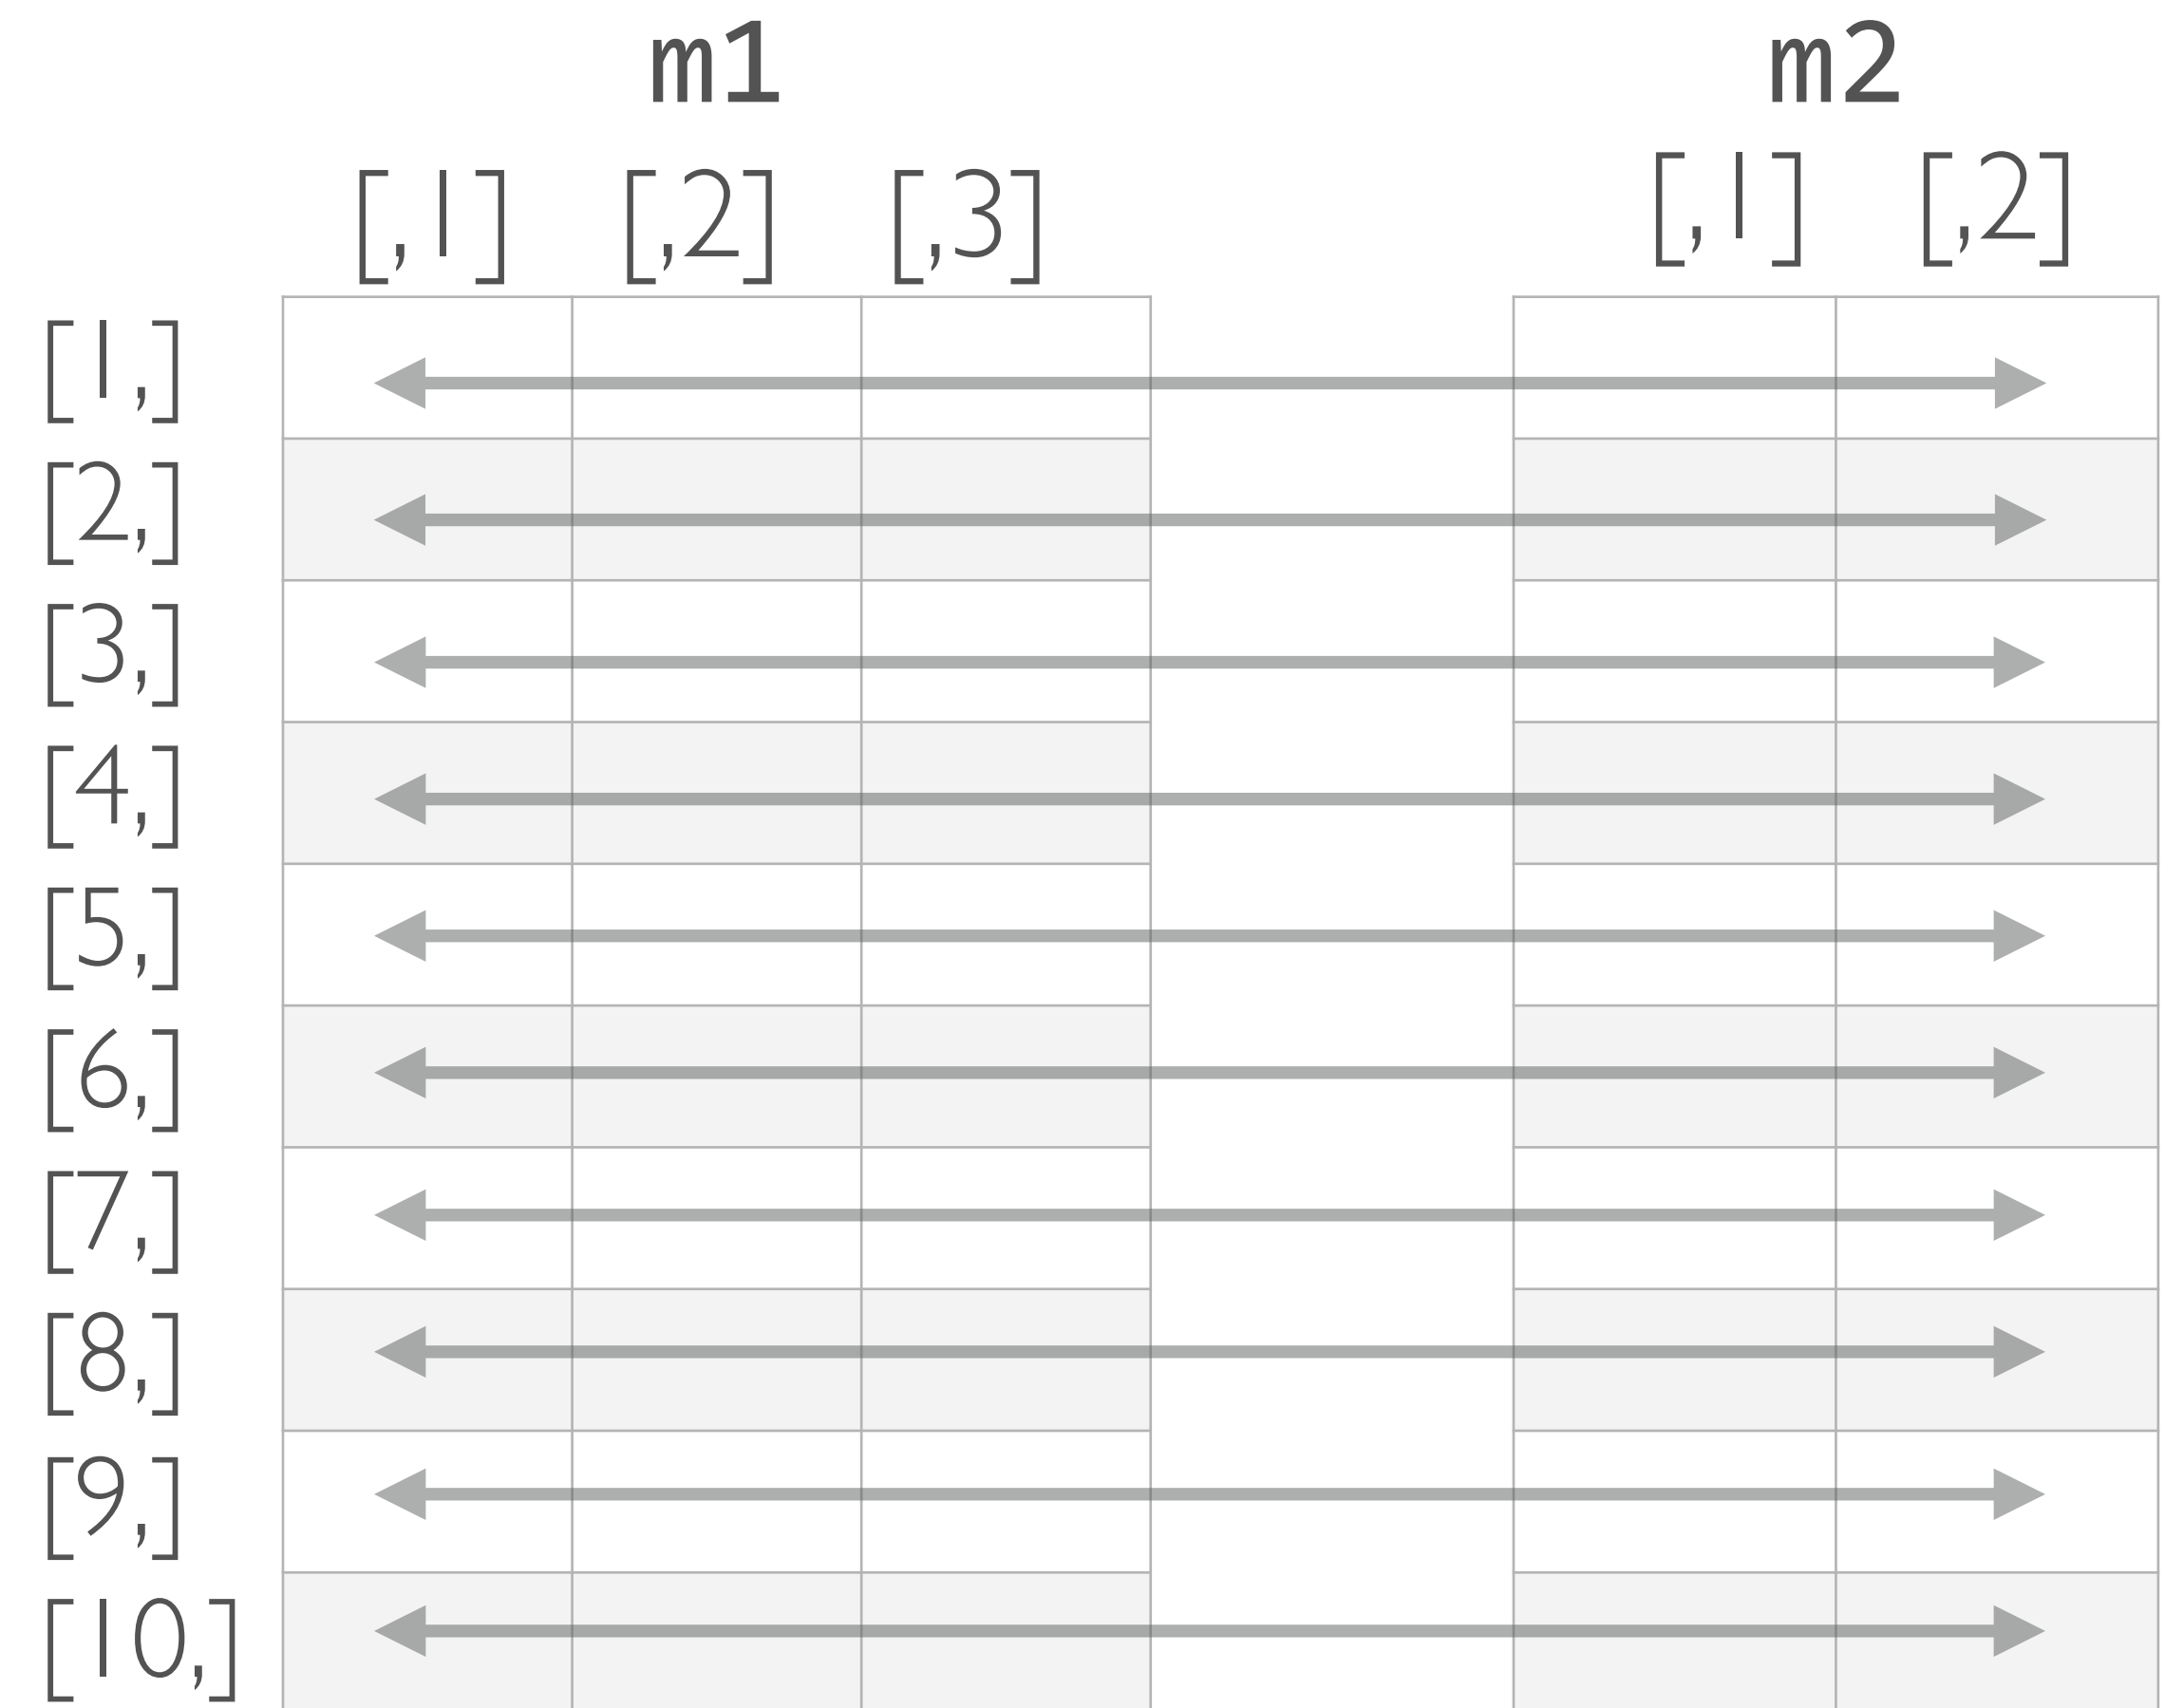
	[,1]	[,2]	[,3]	[,4]	[,5]
[1,]	3	1	10	9	9
[2,]	8	6	5	3	7
[3,]	5	9	7	1	7
[4,]	9	6	6	4	10
[5,]	10	5	2	10	7

*Create this matrix*

# ADDING ON TO



`rbind(m1, m2)`



`cbind(m1, m2)`

# ATTRIBUTES

Matrices have similar attributes as data frames

```
# test these out on your matrix  
dim(m1)  
length(m1)  
str(m1)  
colnames(m1) <- paste("col", 1:5)  
rownames(m1) <- paste("row", 1:5)
```



# INDEXING/SUBSETTING

`matrix[row, col]`

Try these different forms of indexing & subsetting:

```
# extract individual elements
```

```
m1[1, 3]
```

```
m1["row 4", "col 3"]
```

```
# extract all rows and columns 1 through 3
```

```
m1[, 1:3]
```

```
m1[, c("col 1", "col 2", "col 3")]
```

```
# index for all rows and just the second column
```

```
m1[, 2]
```

```
m1[, 2, drop = FALSE]
```

# QUICK SUMMARIES

Get a quick summary of your matrix with **summary()** or any other math/logical operation:

```
summary(m1)
mean(m1)
mean(m[1,])
rowMeans(m1)
colMeans(m1)
rowSums(m1)
colSums(m1)
m > .5
sum(m > .5)
which(m > .5)
m[m > .5]
```

*These same functions can be applied to data frames / tibbles*

# YOUR TURN!

Using the built-in **VADeaths** matrix data:

- 1. Calculate averages for each column and row
- 2. Can you figure out how to add these averages to your table so the output looks like:

	Rural	Male	Rural	Female	Urban	Male	Urban	Female	Avg_by_Age
50-54		11.70		8.70		15.40		8.40	11.050
55-59		18.10		11.70		24.30		13.60	16.925
60-64		26.90		20.30		37.00		19.30	25.875
65-69		41.00		30.90		54.60		35.10	40.400
70-74		66.00		54.30		71.10		50.00	60.350
Avg_by_Local		32.74		25.18		40.48		25.28	30.920

# SOLUTION

```
# Calculate average for each age group and add as a new column
```

```
Avg_by_Age <- rowMeans(VADeaths)
```

```
VADeaths <- cbind(VADeaths, Avg_by_Age)
```

```
# Calculate average for each column and add as a new row
```

```
Avg_by_Local <- colMeans(VADeaths)
```

```
VADeaths <- rbind(VADeaths, Avg_by_Local)
```

VADeaths

	Rural Male	Rural Female	Urban Male	Urban Female	Avg_by_Age
50-54	11.70	8.70	15.40	8.40	11.050
55-59	18.10	11.70	24.30	13.60	16.925
60-64	26.90	20.30	37.00	19.30	25.875
65-69	41.00	30.90	54.60	35.10	40.400
70-74	66.00	54.30	71.10	50.00	60.350
Avg_by_Local	32.74	25.18	40.48	25.28	30.920

# VECTORS

```
[1] 0.67149785 0.47398715 0.32813279 0.87295142 0.56274062 0.16796701 0.05765868 0.59618446  
[9] 0.94417744 0.83129550 0.38959025 0.99178460
```

# PROPERTIES

- 1 dimension
- Can only contain *homogenous* data

```
[1] "a" "b" "c" "d" "e" "f" "g" "h" "i" "j" "k" "l" "m"  
[14] "n" "o" "p" "q" "r" "s" "t" "u" "v" "w" "x" "y" "z"
```

```
[1] 1 2 3 4 5 6 7 8 9 10 11 12 13 14 15 16 17  
[18] 18
```

```
[1] TRUE FALSE TRUE TRUE TRUE FALSE TRUE FALSE  
[9] TRUE TRUE TRUE TRUE FALSE TRUE
```

# CREATING

- Most common way to create a vector is with `c()` or :
- For numeric vectors there are numerous ways to generate sequences of numbers

```
# vectors with no set sequence
c("Learning", "to", "create", "character", "vectors")
c(3, 2, 10, 55)
c(TRUE, FALSE, FALSE, FALSE, TRUE)
```

```
# numeric vectors with regular sequence
6:15
15.5:-6.75
seq(from = 5, to = 95, by = 15)
seq(from = 5, to = 95, length = 4)
```

```
# regular sequence for any type of vector
rep(c(TRUE, TRUE, FALSE), times = 3)
rep(c(TRUE, TRUE, FALSE), each = 3)
```

# CREATING

- Most common way to create a vector is with `c()` or :
- For numeric vectors there are numerous ways to generate sequences of numbers

```
# vectors with no set sequence
c("Learning", "to", "create", "character", "vectors")
c(3, 2, 10, 55)
c(TRUE, FALSE, FALSE, FALSE, TRUE)

# numeric vectors with regular sequence
6:15
15.5:-6.75
seq(from = 5, to = 95, by = 15)
seq(from = 5, to = 95, length = 4)

# regular sequence for any type of vector
rep(c(TRUE, TRUE, FALSE), times = 3)
rep(c(TRUE, TRUE, FALSE), each = 3)
```

There are also many distribution functions to generate data:

- uniform: `<r,d,p,q>unif`
- normal: `<r,d,p,q>norm`
- binomial: `<r,d,p,q>binom`
- poisson: `<r,d,p,q>pois`
- exponential: `<r,d,p,q>exp`



# ADDING ON TO

- Most common way to create a vector is with `c()`
- Combining two different kinds of vectors will coerce the vector to the “simplest” form

```
v1 <- 1:10
v2 <- c(12, 15)
v3 <- c(20, 25:30)
c(v1, v2, v3)
[1] 1 2 3 4 5 6 7 8 9 10 12 15 20 25 26 27 28 29 30
```

```
v4 <- c("Counting from")
c(v4, v1)
[1] "Counting from" "1"          "2"          "3"          "4"
[6] "5"             "6"          "7"          "8"          "9"
[11] "10"
```

```
paste(v4, v1)
[1] "Counting from 1" "Counting from 2" "Counting from 3" "Counting from 4"
[5] "Counting from 5" "Counting from 6" "Counting from 7" "Counting from 8"
[9] "Counting from 9" "Counting from 10"
```

# ATTRIBUTES

Vectors have limited attributes

```
# create this vector
v1 <- 1:10

# try these out on your vector
length(v1)
str(v1)

names(v1)
names(v1) <- paste("Var", LETTERS[1:10])
names(v1)
v1
```

# INDEXING / SUBSETTING

`vector[element]`

Try these different forms of indexing & subsetting:

```
v1[4]  
v1[4:7]  
v1[c(4, 3, 4)]  
v1[c("Var A", "Var D", "Var J")]  
v1[v1 > 6]  
v1[v1 > 8 | v1 <=3]
```

# QUICK SUMMARIES

Get a quick summary of your vector with **summary()** or any other math/logical operation:

```
summary(v1)
```

```
mean(v1)
```

```
median(v1[c("Var A", "Var D", "Var J")])
```

```
v1 > 5
```

```
sum(v1 > 5)
```

# YOUR TURN!

1. check out the built-in character vector *state.name*
2. how many elements are in this vector
3. Can you name each vector element with “V1”, “V2”, ..., “V50”?
4. Subset state.name for those elements with the following names: V35, V17, V14, V38

# SOLUTION

```
# check out state.name  
state.name
```

```
# how many elements are in state.name  
length(state.name)
```

```
# name state.name with “V1”, “V2”,..., “V50”  
names(state.name) <- paste0("V", 1:50)
```

```
# subset state.name for V35, V17, V14, V38  
state.name[c("V35", "V17", "V14", "V38")]
```

# LISTS

```
$item1  
[1] 1 5 3 7
```

```
$item2  
[1] "g" "b" "q" "v" "d" "z" "w" "i"
```

```
$item3  
      [,1] [,2] [,3]  
[1,]    1    4    7  
[2,]    2    5    8  
[3,]    3    6    9
```

```
$item4  
      mpg cyl  disp  hp  drat    wt  qsec vs am gear carb  
Mazda RX4           21.0   6  160  110  3.90 2.620 16.46  0  1    4    4  
Mazda RX4 Wag       21.0   6  160  110  3.90 2.875 17.02  0  1    4    4  
Datsun 710          22.8   4  108   93  3.85 2.320 18.61  1  1    4    1  
Hornet 4 Drive      21.4   6  258  110  3.08 3.215 19.44  1  0    3    1  
Hornet Sportabout  18.7   8  360  175  3.15 3.440 17.02  0  0    3    2  
Valiant             18.1   6  225  105  2.76 3.460 20.22  1  0    3    1
```

# PROPERTIES

- 1 dimension
- Can only contain heterogeneous data - to include multiple and different objects (i.e. vectors, data frames, matrices, and even lists)

```
$item1  
[1] 1 5 3 7
```

```
$item2  
[1] "g" "b" "q" "v" "d" "z" "w" "i"
```

```
$item3  
      [,1] [,2] [,3]  
[1,]    1    4    7  
[2,]    2    5    8  
[3,]    3    6    9
```

```
$item4  
      mpg cyl  disp  hp drat   wt  qsec vs am gear carb  
Mazda RX4           21.0   6  160 110 3.90 2.620 16.46  0  1    4    4  
Mazda RX4 Wag       21.0   6  160 110 3.90 2.875 17.02  0  1    4    4  
Datsun 710           22.8   4  108  93 3.85 2.320 18.61  1  1    4    1  
Hornet 4 Drive       21.4   6  258 110 3.08 3.215 19.44  1  0    3    1  
Hornet Sportabout   18.7   8  360 175 3.15 3.440 17.02  0  0    3    2  
Valiant             18.1   6  225 105 2.76 3.460 20.22  1  0    3    1
```

*Lists are very important objects in R!*  
*They may be confusing but they are worth learning*



# CREATING

- To create a list we use `list()`

```
# list of 4 items
l1 <- list(item1 = 1:3,
           item2 = letters[1:5],
           item3 = c(T, F, T, T),
           item4 = matrix(1:9, nrow = 3))
```

```
l1
## $item1
## [1] 1 2 3
##
## $item2
## [1] "a" "b" "c" "d" "e"
##
## $item3
## [1] TRUE FALSE TRUE TRUE
##
## $item4
##      [,1] [,2] [,3]
## [1,]    1    4    7
## [2,]    2    5    8
## [3,]    3    6    9
```

# ADDING ON TO

- We can add on to lists a couple different ways

```
# add a 5th (named) list item
l1$item5 <- flights
```

```
l1
$item1
[1] 1 2 3
```

```
$item2
[1] "a" "b" "c" "d" "e"
```

```
$item3
[1] TRUE FALSE TRUE TRUE
```

```
$item4
      [,1] [,2] [,3]
[1,]    1    4    7
[2,]    2    5    8
[3,]    3    6    9
```

```
$item5
# A tibble: 336,776 × 19
   year month   day dep_time sched_dep_time dep_delay arr_time sched_arr_time arr_delay
  <int> <int> <int>   <int>         <int>         <dbl>   <int>         <int>         <dbl>
1  2013     1     1     517             515           2     830             819           11
2  2013     1     1     533             529           4     850             830           20
```

# UNDERSTANDING YOUR LIST

Lists have a few main attributes that you care about:

```
str(l1)
```

```
length(l1)
```

```
names(l1)
```

# INDEXING / SUBSETTING

- Its important that you know how to index/subset a list
- Elements of lists can be extracted using 3 approaches:

preserve: `list[component]`

simplify: `list[[component]]`

simplify: `list$component`

```
# try these on our l1 list
l1["item5"]
l1[["item5"]]
l1$item5
l1[["item5"]][1:20, 1:5]
```

*How do they differ?*

# WHAT YOU NEED TO KNOW

- Many statistical modeling results come in the form of lists
- You need to know how to extract parts of a list to access model results

# WHAT YOU NEED TO KNOW

- Many statistical modeling results come in the form of lists
- You need to know how to extract parts of a list to access model results

```
# here's a linear regression model
model <- lm(mpg ~ wt, data = mtcars)

summary(model)
##
## Call:
## lm(formula = mpg ~ wt, data = mtcars)
##
## Residuals:
##      Min       1Q   Median       3Q      Max
## -4.5432 -2.3647 -0.1252  1.4096  6.8727
##
## Coefficients:
##              Estimate Std. Error t value Pr(>|t|)
## (Intercept)  37.2851     1.8776   19.858  < 2e-16 ***
## wt          -5.3445     0.5591   -9.559  1.29e-10 ***
##
```

# WHAT YOU NEED TO KNOW

- Model is simply a list of statistical results for our regression model

```
# here's a linear regression model
model <- lm(mpg ~ wt, data = mtcars)
```

```
names(model)
```

```
## [1] "coefficients" "residuals"      "effects"        "rank"
## [5] "fitted.values" "assign"          "qr"             "df.residual"
## [9] "xlevels"      "call"           "terms"          "model"
```

```
str(model)
```

```
## List of 12
## $ coefficients : Named num [1:2] 37.29 -5.34
##   ..- attr(*, "names")= chr [1:2] "(Intercept)" "wt"
## $ residuals      : Named num [1:32] -2.28 -0.92 -2.09 1.3 -0.2 ...
##   ..- attr(*, "names")= chr [1:32] "Mazda RX4" "Mazda RX4 Wag" "Datsun 710" "Hornet 4 Drive" ...
## $ effects        : Named num [1:32] -113.65 -29.116 -1.661 1.631 0.111 ...
##   ..- attr(*, "names")= chr [1:32] "(Intercept)" "wt" "" "" ...
## $ rank           : int 2
## $ fitted.values: Named num [1:32] 23.3 21.9 24.9 20.1 18.9 ...
```

# WHAT YOU NEED TO KNOW

- Model is simply a list of statistical results for our regression model
- So if you want to extract the residuals or fitted values you can just use normal list subsetting procedures

```
# extract the regression model residuals
model$residuals
##           Mazda RX4           Mazda RX4 Wag           Datsun 710
##          -2.2826106          -0.9197704          -2.0859521
##      Hornet 4 Drive      Hornet Sportabout           Valiant
##           1.2973499          -0.2001440          -0.6932545
##           Duster 360           Merc 240D           Merc 230
##          -3.9053627           4.1637381           2.3499593
##           Merc 280           Merc 280C           Merc 450SE
##           0.2998560          -1.1001440           0.8668731
##           Merc 450SL           Merc 450SLC Cadillac Fleetwood
##          -0.0502472          -1.8830236           1.1733496
## Lincoln Continental Chrysler Imperial           Fiat 128
##           2.1032876           5.9810744           6.8727113
```



# YOUR TURN!

1. *Create this regression model:*

```
flight_lm <- lm(arr_delay ~ dep_delay + month + carrier,  
               data = flights)
```

2. Extract the residuals from the **flight\_lm** list

3. What is the min, max, median, and mean of these residuals?

# SOLUTION

```
# create regression model
flight_lm <- lm(arr_delay ~ dep_delay + month + carrier, data = flights)

# extract residuals from flight_lm list
residuals <- flight_lm$residuals

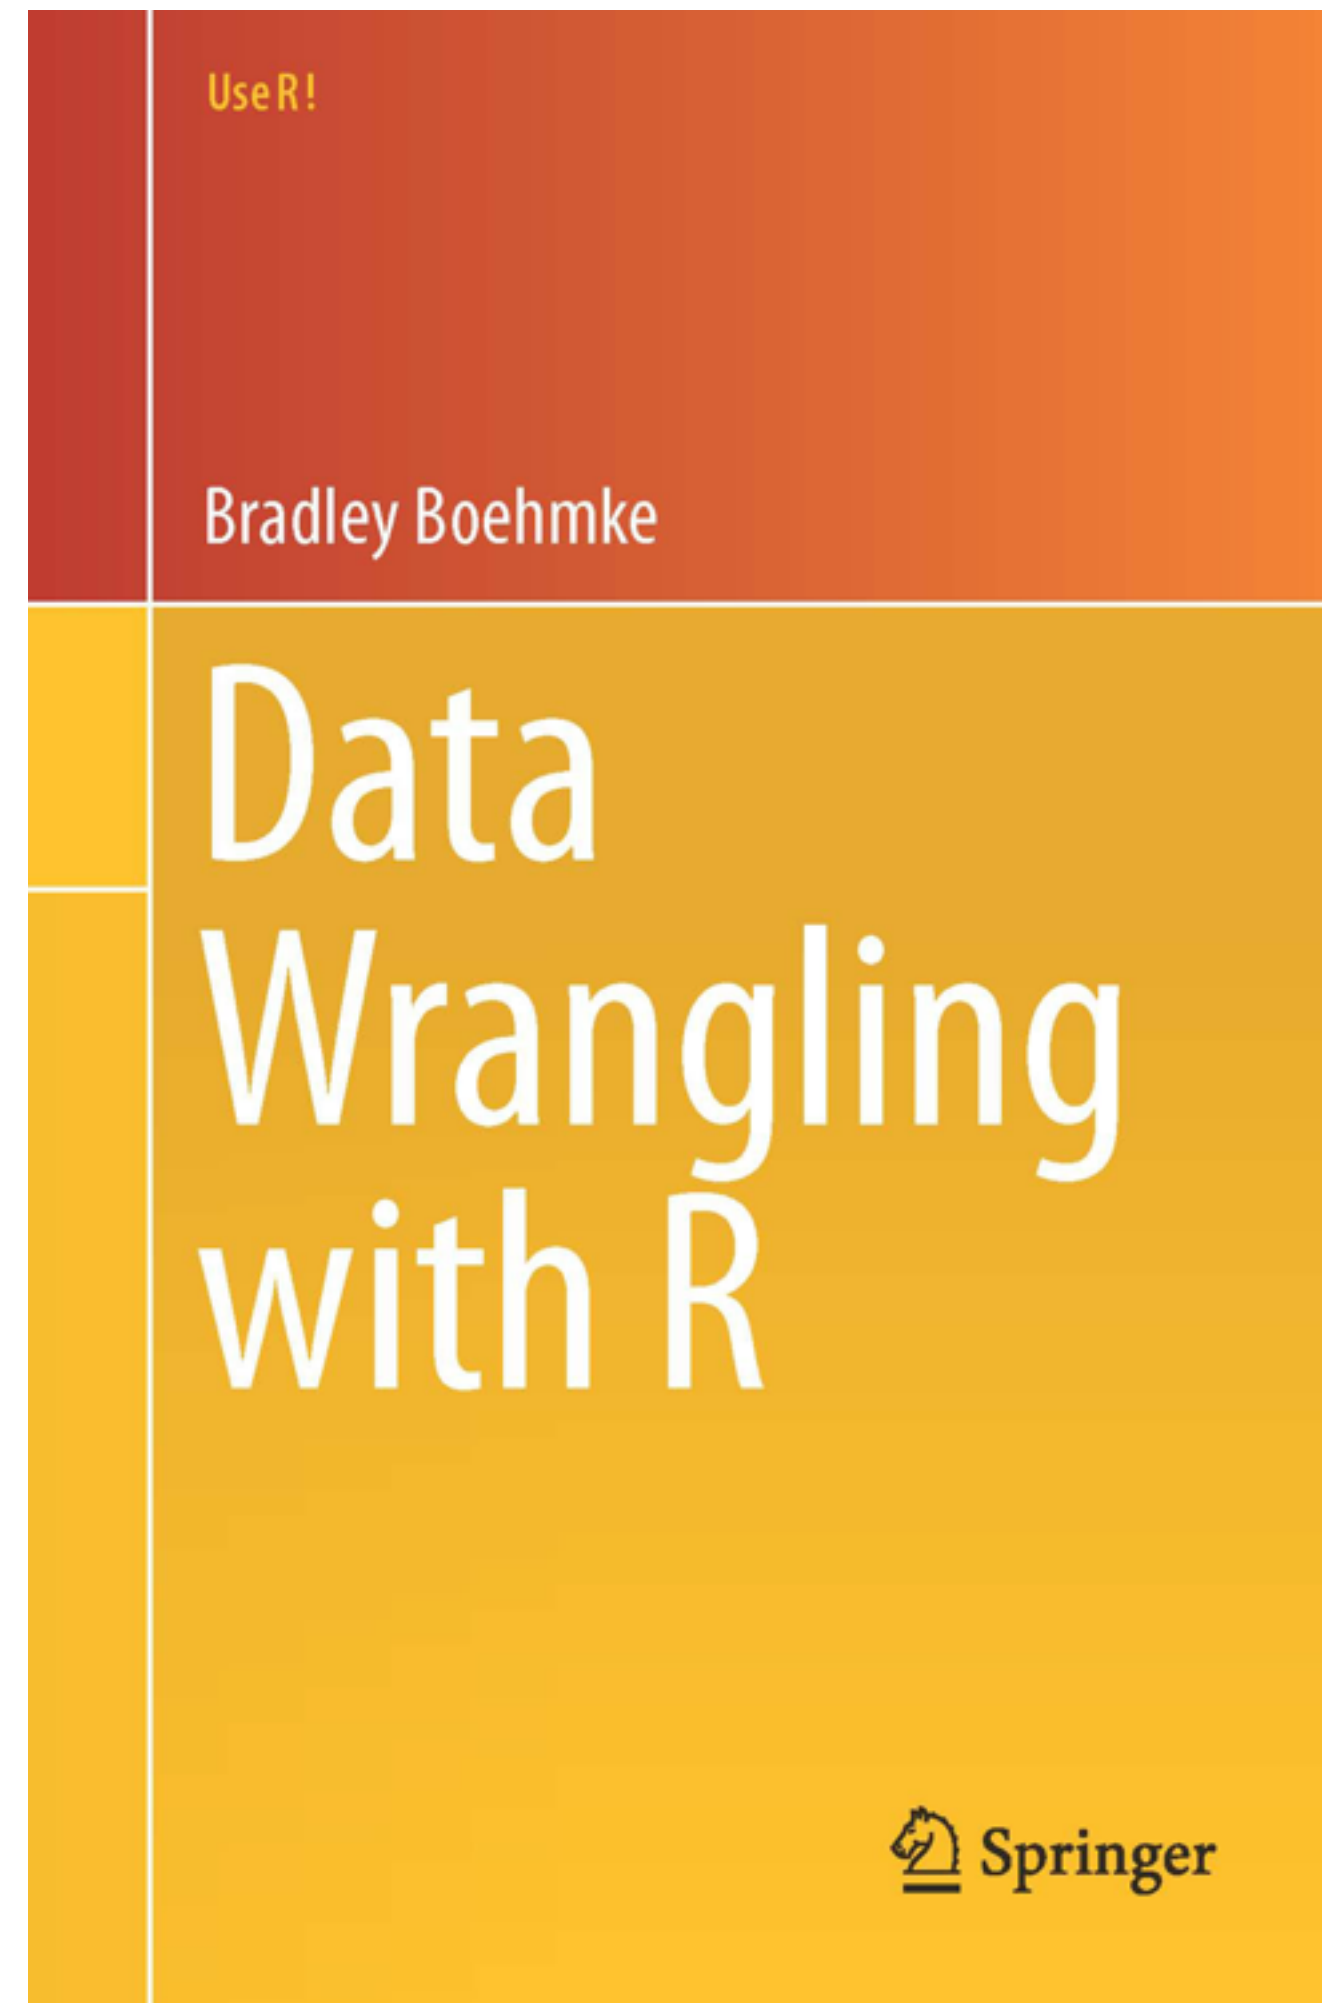
# compute summary statistics
summary(residuals)
```

SO LITTLE TIME!





# LEARN MORE



WHAT TO REMEMBER



# FUNCTIONS TO REMEMBER

Operator/Function	Description
<code>data.frame</code> , <code>as_tibble</code> , <code>matrix</code> , <code>list</code> , <code>c()</code> , <code>:</code>	create data frames, tibbles, matrices, etc.
<code>str</code> , <code>names</code> , <code>colnames</code> , <code>rownames</code> , <code>dim</code> , <code>length</code> , <code>nrow</code> , <code>ncol</code>	understand attributes of data structures
<code>summary</code> , <code>mean</code> , <code>median</code> , <code>sum</code> , <code>colSums</code> , <code>rowSums</code> , <code>colMeans</code> , <code>rowMeans</code>	understand summary statistics of data structures
<code>[]</code> , <code>[[[]]</code> , <code>\$</code>	index & subset data structures